

THE MODEL PICKER

Which AI for which job

Three rules that survive every launch week — and
the models filling each slot.



usd per million tokens · input / output

RULE 1

A flagship for thinking

Proposals, contracts, hard decisions — the work where being wrong is expensive.



Claude Opus 5.5

at or near the top of the main rankings

\$4 / \$20



GPT-6 Astra

more than twice the price

\$10 / \$50

RULE 2

A cheap, fast model for sorting

Tagging emails, summarising calls, reading invoices — thousands of times a month.



GPT-6.1 Sol

a fifth of Astra's price

\$2 / \$10



Gemini 3.8 Flash

cheapest of the three

\$0.75 / \$3.75


UNTIL 31 DEC 2026

WHAT THAT MEANS IN MONEY

Sorting 3,000 emails a month


Same job. Same volume. Three bills.

GPT-6 Astra



~\$30

GPT-6.1 Sol



~\$6

Gemini 3.8 Flash



~\$2

Rough estimate, about 1,000 tokens per email.

RULE 3

Use the AI that lives where your work is

A model you have to leave your workflow to reach is a model you stop using by Thursday.



Gemini

Gmail · Docs



Copilot

Microsoft 365

Not the flagship

Released 28 September. Within about 2 points of Opus 5.5 on coding tests, at half the price.



Claude Sonnet 5.5

the coding default

\$2 / \$10

Here the flagship earns it

It writes the animation as code and renders the finished video. Designers call motion its biggest jump.



Claude Opus 5.5

writes the animation, renders the video

\$4 / \$20

A speech model, not a text one

One that listens and talks in a single step, instead of transcribing, thinking, then speaking.

~0.5s

where replies now land

THE ACTUAL STRATEGY

Mix them

Nobody picks one model. You build a chain.

01 Everything arrives
emails, calls, invoices, messages



02 The cheap model sorts
tags, summarises, throws out the noise



03 The flagship sees what's left
only the work that needs judgment

Names and prices change every month. **These rules don't.**

**Which AI subscription are
you paying for that you
barely use?**

@shazimbuid